

CS/COE 0447

Introduction to:
Introduction to
Computer Architecture

wilkie (with content borrowed from:
Jarrett Billingsley
Dr. Bruce Childers)

Memory

Where data calls “home”

What is the memory?

- the **system memory** is a piece of *temporary* storage hardware
 - it's smaller and faster than the *persistent* storage. (disk, etc)
 - maybe in the future it won't be temporary, or the line between system memory and persistent storage will go away...
- it's where the **programs and data** that the computer is currently executing and using reside
 - all the variables, all the functions, all the open files etc.
 - the CPU **can only run programs from system memory!**

Memory: Looking at Bytes

- the memory is a **big one-dimensional array of bytes**
- what do these bytes mean?
 - `_(ツ)_/`
- every byte value has an **address**
 - this is its "array index"
 - addresses start at 0, like arrays in C/Java
 - gee wonder where they got the idea
- when each byte has its own address, we call it a **byte-addressable machine**
 - not many *non*-byte-addressable machines these days
 - A non-byte-addressable machine would have addresses that refer to **words**, etc. (Getting data in-between words may be difficult... it is called misaligned data. It's a yuck problem.)
 - Thankfully, we are using a byte-addressable architecture

Add r	Val
0	00
1	30
2	04
3	00
4	DE
5	C0
6	EF
7	BE
8	6C
9	34
A	00
B	01

Memory: It's Finite (but how much?)

- each address refers to *one* byte. if your addresses are n bits long... **how many bytes** can your memory have?
 - 2^n B
- machines with 32-bit addresses can access $2^{32} \text{ B} = \mathbf{4GiB}$ of memory
 - with 64-bit addresses... **16EiB** lol
- kibi, mebi, gibi, tebi, pebi, exbi are powers of 2
 - **kiB** = 2^{10} , **MiB** = 2^{20} , **GiB** = 2^{30} etc.
- kilo, mega, giga, tera, peta, exa are *ostensibly* powers of 10
 - **kB** = 10^3 , **MB** = 10^6 , **GB** = 10^9 etc.
- **but most people still say "kilo, mega" to mean the powers of 2**
 - really only hard drive manufacturers use the "power of 10"
 - 1TB hard drive is 10^{12} B ... $10^{12} \div 2^{40} = \mathbf{909 \text{ GiB}}$
 - now you know why it's like that

Memory: Looking at Words

- for most things, we want to use **words**
 - the "comfortable" integer size for the CPU
 - on this version of MIPS, it's **32b (4B)**
- but our memory only holds bytes... wat do
- **combine multiple bytes into larger values**
 - the CPU can handle this for us
 - but importantly, *the data is still just bytes*
- when we talk about values bigger than a byte...
 - the **address** is **the address of their first byte**
 - the byte at the *smallest* address
 - so what are the addresses of the three words here?

Addr	Val
0	00
1	30
2	04
3	00
4	DE
5	C0
6	EF
7	BE
8	6C
9	34
A	00
B	01
C	02

Endianness

Because We Like to Make Numbers More Confusing, Honestly

A Matter of Perspective

- Let's say there's a word at address 4... made of 4 bytes (32-bits)
- *Wh...what word do those 4 bytes represent?*

If we think
of addresses
*increasing
downward...*
...is it
0xDEC0EFBE?

Addr	Val
...	...
4	DE
5	C0
6	EF
7	BE
...	...

Addr	Val
...	...
7	BE
6	EF
5	C0
4	DE
...	...

if we think of
addresses
*increasing
upward...*
...is it
0xBEEFC0DE?

The Two Endians

- When interpreting a *sequence of bytes* as larger values, **endianness** is the rule used to decide **what order to put the bytes in**

big-endian means the **first byte** is the "big end:" the **most significant** byte

0x**DE**C0EFBE

0	1	2	3
DE	C0	EF	BE

little-endian means the **first byte** is the "little end:" the **least significant** byte

0xBEEFC0**DE**

first byte = **byte at smallest address**
nothing to do with **value** of bytes, only **order**

Which is “better” ?

- it doesn't matter.* **as long as you're consistent**, it's fine
- for political _(x86) reasons, most computers today are little-endian
- but endianness pops up whenever you have sequences of bytes:
 - like in files
 - or networks
 - or hardware buses
 - or... memory!
- which one is MIPS?
 - it's *bi-endian*, meaning it can be configured to work either way
 - **but** MARS uses the **endianness of the computer it's running on**
 - so **little-endian** for virtually everyone
 - cause x86
 - *ugh*

What Doesn't Endianness Affect?

- × the arrangement of the bits *within a byte*
 - it just changes **meaning of order of the bytes**
 - note the bytes are still **DE**, **C0** etc.
- × **1-byte** values, arrays of bytes, **ASCII** strings...
 - *single bytes* don't care about endianness at all
- × the ordering of bytes **inside the CPU**
 - there's no need for e.g. "big-endian" arithmetic
 - the CPU works with whole *words*
- endianness only comes up when moving data:
 - larger than single bytes
 - between **the CPU and memory**
 - or between **multiple computers**

0xBEEFC0DE

0xED0CFEEB


0xDEC0EFBE

l l e H o

H e l l o

The Messy Origin of Endianness

- There was once a nation called Lilliput.
- Once, a future king cut themselves when breaking their boiled egg big-end first.
- So the emperor outlawed such a practice and eggs can only be cut my their little end first.
- Many rebellions, civil wars, and executions happened on both sides as each fought to win the ability to eat their eggs whichever way. The big-endians vs. the little-endians.
- It's war/political satire from "Gulliver's Travels" by Jonathan Swift. Used again as satire to ***complain about how ridiculous computer endianness is.*** *~~ Infinite sighs ~~*



Variables, Loads, and Stores

What are loads? Stores? And why Variables are Not Scary in asm.

Memory Addresses

- *everything* in memory has an **address**
 - the position in memory **where it begins**
 - where its **first byte** is
 - this applies to variables, functions, objects, arrays etc.
- a super important concept:

**every variable really has two parts:
an **address** and a **value****

- if you want to put a variable in memory...
 - first you need to figure out **what address to put it in**
 - this *extremely tedious* task is handled by assemblers
 - whew

Declaring Global Variables

- we can declare a global variable like this (at the top of a file):


.data

x: .word 4

- the Java equivalent would be **static int x = 4;**
- **.data** says "I'm gonna declare variables"
 - you can declare as many as you want!
 - to go back to writing code, use **.text**
- if we assemble this little program and make sure **Tools > Show Labels Window** is checked, what do you see?
 - the assembler *gave the variable* that address
 - it'll do that for every variable

Load-Store Architectures

- in some architectures, *many* instructions can access memory
 - x86-64: **add** [rsp-8], rcx
 - adds the contents of **rcx** to the value at address **rsp-8**
- in a **load-store** architecture, **all** memory accesses are done with two kinds of instructions: loads and stores (like in MIPS)




loads copy data **from** memory
into CPU registers

stores copy data **from** CPU
registers **into** memory

Operating On Variables in Memory

- we want to increment a variable that is **in memory**
 - where do values **have to be for the CPU to operate on them?**
 - what do we want the **overall outcome to be?**
- so, what **three steps** are needed to increment that variable?
 1. **load** the value from memory into a register
 2. **add 1** to the value in the register
 3. **store** the value back into memory
- **every variable access** works like this!!!
 - High Level Languages (HLLs) just hide this from you



Accessing Memory in MIPS


MIPS ISA: load/store words

- You can load and store entire 32-bit words with **lw** and **sw**
- The instructions look like this (variable names not important):

lw t1, x # loads *from* variable x *into* t1

sw t1, x # stores *from* t1 *into* variable x

- In MIPS, **stores are written with the destination on the *right*. !?**
 - Well, you can remember it with this diagram...
 - The memory is "on the right" for both loads and stores



Read, Modify, Write

- you now know enough to increment x!

- first we **load x into a register**

```
lw    t0, x
```

- then...

```
add   t0, t0, 1
```

- and then...

```
sw    t0, x
```

- Let's see what values are in **t0** and memory after this program runs

Variables: That's really it.

- Variables in asm aren't *THAT* scary
- *Please* don't be afraid of them
- You just gotta remember to store if you wanna change em
- And remember... you are the CPU:
 - Loads: Read from memory
 - Stores: Write to memory

Smaller Values

When your 32-bit cup doth not overfloweth... wait, no, that's not right.

MIPS ISA: load/store bytes/half-words

- some values are `tiny`
- to load/store **bytes**, we use **lb/sb**
- to load/store 16-bit (**half-word**) values, we use **lh/sh**
- These mostly look and work just like **lw/sw**, like:

lb `t0, tiny` # loads a byte into `t0`

sb `t0, tiny` # stores a byte into `tiny`

- I said mostly... recall: how big are registers?
 - So, what should go in those extra 16/24 bits then?
 - ???

Can I Get an Extension?

- Sometimes you need to *widen* a number with fewer bits to more
- **zero extension** is easy: **put 0s at the beginning.**

$1001_2 \rightarrow \text{to 8 bits} \rightarrow 0000\ 1001_2$

- But there are also **signed numbers** which we didn't talk about yet... hmm

Signed Numbers (sign-magnitude)

- Seems like a good time to think about “negative” values.
 - These are numbers that have nothing good to say.
- Binary numbers have bits which are either 0 or 1.
 - Well, yeah...
- So what if we used one bit to designate “positive” or “negative”
 - Called **sign-magnitude** encoding:

$$\underbrace{10100010}_{\text{sign-magnitude}} = \underbrace{-34}$$

$$\underbrace{00010110}_{\text{sign-magnitude}} = \underbrace{22}_{\text{(normal)}}$$

Signed Numbers (problems)

$$\begin{array}{lcl} \textcolor{red}{1}\underbrace{00000000} & = & \textcolor{red}{-}\underbrace{0} \\ \textcolor{red}{0}\underbrace{00000000} & = & \underbrace{0} \end{array}$$

- Waaaaait a second.
 - What is negative zero???
- This encoding allows two different zeros.
 - This means we can represent how many different values (8-bit)?
 - $2^8 - 1$ (minus the one redundant value) = 255 (-127 ... 0 ... 127)
- Sign-magnitude is a little naïve... let's try a different approach...

Signed Numbers (1's Complement)

- Let's borrow a technique from accounting and mechanical calculators: **flip the dang bits**.

$$11010100 = -\underbrace{00101011} = -\underbrace{43}$$

$$\underbrace{00100110} = \underbrace{00100110} = 38$$


$$\underbrace{00000000} = \underbrace{00000000} = 0$$

$$11111111 = -\underbrace{00000000} = -\underbrace{0} \text{ 😡}$$


- OH COME ON (actually, this *is* better because math is easier)
 - But this is really *isn't* used that much.

Signed Numbers (2's Complement)

- This one, I promise, is juuuuust right.
 - But it's a little strange!
- We'll just make SURE there is only one zero:



1's Complement



2's Complement



- So, we flip the bits... (1's complement) and add one.
 - Adding one makes sure our -0 is used for -1 instead!
- Sure, it's a little lopsided, but, hey, we get an extra number.
 - But, hmm, but -4 **doesn't have a valid positive number**.
 - That's the trade-off, but it's for the best.

Signed Numbers (2's Complement)

- Let's look at the **same bit patterns** as before:

$$11010100 = -\underbrace{00101011} = -(\underbrace{43+1}) = -44$$

$$\underbrace{00100110} = \underbrace{00100110} = 38$$

$$\underbrace{00000000} = \underbrace{00000000} = 0$$

$$11111111 = -\underbrace{00000000} = -(\underbrace{0+1}) = -1$$

- If the MSB is 1: Flip! Add one!
- Otherwise:** Do nothing! It's the same!

Signed Numbers (2's Complement)

- What happens when we add zeros to a positive number:

$$00100110 = 38$$

$$00000000000100110 = 38$$

- What happens when we add ones to a negative number:

$$10100110 = -90$$

$$1111111110100110 =$$

$$-000000000001011001 = -90$$

Dang that's cool!

Can I Get an Extension? (Reprise)

- Sometimes you need to *widen* a number with fewer bits to more

- **zero extension** is easy: **put 0s at the beginning.**

$1001_2 \rightarrow \text{to 8 bits} \rightarrow 0000\ 1001_2$

- But there are also **signed numbers** which we didn't talk about yet

- The **top bit (MSB)** of signed numbers determines **the sign (+/-)**


- **sign extension** puts ***copies of the sign bit*** at the beginning

$1001_2 \rightarrow \text{to 8 bits} \rightarrow 1111\ 1001_2$

$0010_2 \rightarrow \text{to 8 bits} \rightarrow 0000\ 0010_2$


EXPAND

- If you load a **byte**...



10010000

If the byte is **signed**... what *should* it become?



lb does
sign extension.

If the byte is **unsigned**... what *should* it become?




lbu does zero
extension.

lbu (load byte unsigned) is USUALLY what you want to use!

Truncation

- if we go the other way, **the upper part of the value is cut off.**



- the sign issue doesn't exist when storing, cause we're going from a *larger* number of bits to a *smaller* number
 - therefore, **there are no sbu/shu instructions**