

CS/COE 0447

Functions:
Calling Conventions,
and The Stack

wilkie (with content borrowed from:
Jarrett Billingsley
Dr. Bruce Childers)

Lightning Recap

- We can turn if-else pseudo-code into:

```
if(s0 == 30) {
```

```
    block A
```


```
}
```

```
else {
```

```
    block B
```

```
}
```

```
block C
```



```
beq s0, 30, blockA
```

```
b blockB
```

```
blockA:
```

```
    block A
```

```
b blockC
```

```
blockB:
```

```
    block B
```

```
blockC:
```

```
    block C
```

Our Alternative Approach

- If you see a conditional branch followed by an unconditional one...
...you can **merge them** like so:

invert the condition and
get rid of the first label.

```
beq s0, 30, if  
b else
```

```
bne s0, 30, else
```

if:

block A

b end

else:

block B

end:

block A

b end

else:

block B

end:


Calling Conventions

It's kind of like computer etiquette.

What is a Calling Convention (CS 449)

- It ensures our programs don't trip over their own feet
- It's **how machine-language functions call one another**.
 - How **arguments** are passed.
 - How **return values** are returned.
 - How control **flows into/out of the function**
 - What **contracts** exist between the caller and the callee.
 - **For instance: What registers it will use.**
- Common terms we will be using: Caller v. Callee:


```
void major() {  
    minor();  
}  
      caller
```



```
void minor() {  
    ...  
}  
      callee
```

The Program Counter

- a program's instructions are in memory, so they have **addresses**
- the PC (program counter) holds the address of **the next instruction to run**



btw what pattern do you notice about these addresses?

Branches and the Program Counter

- The “branch” instructions we have already seen actually simply interact with the program counter.

Instruction	Meaning
beq a, b, label	if(a == b) { \$pc = label }
bne a, b, label	if(a != b) { \$pc = label }
blt a, b, label	if(a < b) { \$pc = label }
ble a, b, label	if(a <= b) { \$pc = label }
bgt a, b, label	if(a > b) { \$pc = label }
bge a, b, label	if(a >= b) { \$pc = label }

- Each sets \$pc to the address of whatever label is given.

Control Flow (Reprise)

- When the caller calls a function, where do we go?
- When the callee's code is finished, where do we go?

The diagram illustrates the control flow between two functions: `fork()` and `knife()`. `fork()` is labeled as the 'caller' and `knife()` as the 'callee'. Red arrows show the flow of execution: one arrow points from the call `knife();` inside `fork()` to the start of `knife()`, and another arrow points from the closing brace `}` of `knife()` back to the line `spoon++;` in `fork()`.

```
void fork() {  
    knife();  
    spoon++;  
}  
    caller
```

```
void knife() {  
    spork++;  
    spatula--;  
}  
    callee
```


MIPS ISA: The “jal” Instruction

- We **call** functions with **jal**: jump and link

What address should
go into PC next?

When **func** returns,
where will we go?

PC 0x8004

PC 0x8C30

ra 0x8008

0x8000 **li** a0, 10

0x8004 **jal** func

0x8008 **li** v0, 10

...

This is what **jal** does:

It **jumps** to a new location, and
makes a **link** back to the old one in
the **ra** (return address) register

and this is ALL it does.

func: 0x8C30 **li** v0, 4

...

MIPS ISA: The “jr” Instruction

- We return from functions with **jr**: jump to address in register

Now we're at the end of **func**. **ra** still has the proper return address (*thanks to jal*)

jr ra copies **ra** into **pc**.

PC 0x8C38

ra 0x8008



PC 0x8008

0x8000 **li** a0, 10

0x8004 **jal** func

0x8008 **li** v0, 10

... ..

func: 0x8C30 **li** v0, 4

0x8C34 **syscall**

0x8C38 **jr** ra

and this is ALL it does.

Calling Conventions: Arguments and Return Values

Functions that can now do things!

Passing Arguments

- If we have a function in a higher level language...

```
v0
int gcd(a0 int a, a1 int b) {
    while(a != b) {
        if(a > b)
            a -= b;
        else
            b -= a;
    }
    return a;
}
```

We use particular registers to pass arguments and return values.

We already know how to return. How?

For this, just *put the value you want to return in **v0** before jr ra.*

The “a” and “v” Registers

- **a0-a3** are the **argument registers**
- **v0-v1** are the **return value registers**
 - This is just a **convention**, there's nothing special about them
- To call a function...
 - You **put its arguments in the a registers** before doing a jal
- Once control is inside the callee...
 - The arguments are just "there" in the a registers.
 - Cause they are.
 - They didn't go anywhere...

Let's Write a Function

- Let's write a function like this:

```
int add_nums(int x, int y) {return x + y;}
```

- Inside of our **add_nums** asm function...
 - which register represents **x**?
 - which register represents **y**?
 - which register will hold the sum that we return?

add_nums:

```
add v0, a0, a1
```

```
jr    ra
```

Let's Call That Function

- Now let's make **main** do this:
v0 = add_nums(3, 8)
print(v0);
- How do we set 3 and 8 as the arguments?
- How do we call **add_nums**?
- Afterwards, which register holds the sum?
- So how can we print that value?
- *Why* do syscalls put the number of the syscall in **v0**?
 - Well what do you get when you cross an elephant and a rhino?
 - Hell if I know (it's a bad, confusing decision. Sorry.)

```
li    a0, 3
li    a1, 8
jal   add_nums
move  a0, v0
li    v0, 1
syscall
```

Calling Conventions: Saved and Unsaved Registers

It's about how to trust functions to not step on your toes.

An Innocent Looking Function

- Let's make a variable and a function to change it

.data

counter: .word 0

.text

increment:

la t0, counter

lw t1, (t0)

add t1, t1, 1

sw t1, (t0)

jr ra

then we can call it

main:

jal increment

jal increment

jal increment

Everything's Fine Right?

- Let's write a loop that calls it ten times in a row
- So we need a loop counter ('i' in a for loop)

```
li t0, 0 # our counter
loop_begin:
    jal increment
    add t0, t0, 1
    blt t0, 10, loop_begin
loop_end:
(another way to write a for loop)
```

If we run this, it only increments the variable **once**. ☹

Why? let's put a **breakpoint** on blt and see what it sees.

Another Piece of the Calling Convention Puzzle

- When you call a function, **it's allowed to change some registers**
- But other registers **must be left exactly as they were**

functions are required to put these registers **back the way they were before they were called.**

Saved	Unsaved
s0-s7	v0-v1
sp	a0-a3
ra*	t0-t9

anyone can change these. after you call a function, **they might have totally different values from before you called it.**

*ra is a little weird cause it's kinda "out of sync" with the other saved regs but you DO save and restore it like the others

When You Call a Function...

- After a **jal**, you have no idea what's in these registers.

...

jal increment

...

Unsaved

v0-v1

a0-a3

t0-t9

could be nonsense!
garbage! bogus!

Why It Broke

- If we look at this code again...

```
li t0, 0
loop_begin:
jal increment
add t0, t0, 1
blt t0, 10, loop_begin
loop_end:
```

t0 is our loop counter and everything's fiiiine.

uh oh.

WHAT IS IN t0 NOW??

Instead, this is a great place to use an **s register**.

Fixing Our Function (Abiding by Convention)

- If we use an s register...

```
li s0, 0  
loop_begin:
```

```
jal increment
```

```
add s0, s0, 1  
blt s0, 10, loop_begin
```

```
loop_end:
```

s0 is our loop counter and everything's fiiiine.

uh oh.

oh whew, we used an s register, it's fine.

But s registers aren't magic. **they don't do this automatically... (What about \$ra !?)**

Don't Step on Other's Toes

- So, let's look at this problem: track **PC** and **ra** as we run this code.

	PC	ra		
	0x8000	0x0000	0x8000	jal fork
			0x8004	li v0, 10
		
After jal fork:	0x8020	0x8004		
After jal spoon:	0x8040	0x8024	fork:	0x8020 jal spoon
				0x8024 jr ra
After jr ra:	0x8024	0x8024		...
After jr ra:	0x8024	0x8024		...
After jr ra:	0x8024	0x8024	spoon:	0x8040 jr ra
After jr ra:	0x8024	0x8024		
After jr ra:	0x8024	0x8024		
After jr ra:	0x8024	0x8024		

UHHHHHHHHHHH

What's the Deal


- There's only **one return address register**
- If we call more than **one level deep**, things go horribly wrong...
- Could we put it in another register?
 - Then what about three levels deep? four?
 - We just don't have **enough registers...**
 - **Pro-tip:** *Resist any urge to preserve registers in other registers.*
 - Just use memory.
- So **where** do we put things when we don't have room in registers?
 - Memory. Specifically...

The Stack

It's just memory.


A Busy Desk

- There's a tiny desk that three people have to share
- Person 1 is working at the desk. it's covered in their stuff.
- Person 2 **interrupts them** and needs to do some important work
- **What does person 2 do with the stuff?**
 - Throw it in the trash?
- They **put it somewhere else.**
- When they are done, **they put it back.**
 - According to “etiquette” aka “calling convention” 😊




A Very Busy Desk

- So person 2 is working...
- Then person 2 is interrupted by person 3.
- When person 3 is done, person 2 will come back.
- Where do we put person 2's stuff?
 - **on top of the stack of stuff.**
- The desk represents the **registers**.
- The people are **functions**.
- The stack of stuff is... **the stack**.



What is the Stack? (“sp” Register)

- It's an **area of memory** provided to your program by the OS
 - When your program starts, it's already there.
- The stack is holds **information about function calls**:
 - **The return address** to the caller
 - **Copies of registers** that we want to change
 - **Local variables** that can't fit in registers
- How do we access the stack?
 - Through the **stack pointer (sp) register**
 - This register is initialized for you by the OS too.



Accessing the Stack (animated)

- Let's say **sp** starts at the address **0xF000**
- We want to **push** something on the stack
- The first thing we'll do is **move sp to the next available slot**
- Clearly, that's the *previous* address
 - **Subtract** 4 from **sp** (why 4?)
- Then, we can store something in **the memory that sp points to.**



sp →

...	...
0xF008	0x00000000
0xF004	0x00000000
0xF000	0x00000000
0xEFFC	0xC0DEBEEF

MIPS: Accessing the Stack (animated)

- Say `ra = 0xC0DEBEEF`
- First: move the stack pointer down (up?):

sub `sp, sp, 4`

- Then, store `ra` at the address that `sp` holds.

sw `ra, (sp)`

- Now the value in `ra` is **saved** on the stack, and **we can get it back later**.
 - And we can store as many return addresses as we want!

`sp` →

...	...
0xF008	0x00000000
0xF004	0x00000000
0xF000	0x00000000
0xEFFC	0xC0DEBEEF

MIPS: Going the Other Direction (animated)

- Now we wanna **pop** the value off the stack and put it back in **ra**
- We do the same things, but in reverse

lw **ra**, (**sp**)

- Then, we move the stack pointer...
up? down? whatever

add **sp**, **sp**, **4**

- Now we got back the old value of **ra**!
- And **sp** is back where it was before!

ra **0xC0DEBEEF**

...	...
0xF008	0x00000000
0xF004	0x00000000
0xF000	0x00000000
sp → 0xEEFC	0xC0DEBEEF

MIPS: Simplifying: “push” and “pop” pseudo-ops

- The push and pop operations always look and work the same
- Since you'll be using them in most functions, we shortened em!
- If you write **push ra** or **pop ra**, it'll do these things for you!
 - Thank goodness.

push ra = **subi** sp, sp, 4
 sw ra, (sp)

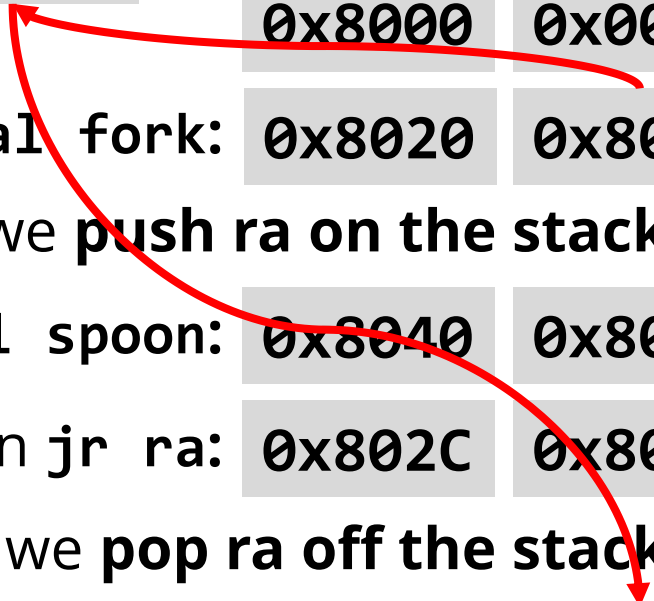
pop ra = **lw** ra, (sp)
 addi sp, sp, 4

These are **pseudo-ops**: fake instructions to shorten common tasks

These can be used with ANY register, not just **ra**!

Some Steel-Toed Boots

sp	➡	0x8004	PC	ra
			0x8000	0x0000
After jal fork:			0x8020	0x8004
Then we push ra on the stack!				
After jal spoon:			0x8040	0x802C
After spoon jr ra:			0x802C	0x802C
Then we pop ra off the stack!				
Before fork jr ra:			0x8034	0x8004
After fork jr ra:			0x8004	0x8004



0x8000	jal	fork
fork:		
0x8020	push	ra
0x8028	jal	spoon
0x802C	pop	ra
0x8034	jr	ra
spoon:		
0x8040	jr	ra

Writing a Simple Function

- Writing a function follows a simple structure:

1. Give it a name (label). **spoon:**

2. Save **ra** to the stack. **push ra**

3. Do whatever. ***your code goes here***

4. Load ra from the stack. **pop ra**

5. Return! **jr ra**

- Push everything you may need. Pop it back at the end.

Extending: Preserving other Registers

- Writing a function follows a simple structure:

1. Give it a name (label). **spoon:**

2a. Save **ra** to the stack.

push ra

2b. Save **s0** to the stack.

push s0

3. Do whatever.

your code using s0 goes here

4a. Load s0 from the stack.

pop s0

4b. Load ra from the stack.

pop ra

5. Return!

jr ra

- Push everything you may need. Pop it back **in reverse order** at the end. **(We will look at why we do this and when soon!)**

When to “push” and “pop”

- Treat pushes and pops like the { **braces** } around a function.

spoon:

```
push ra # {  
# 800 instructions  
# so much stuff omg  
pop  ra # }  
jr   ra
```

pushes come at the
beginnings of functions

pops come at the end

**That is it, seriously, don't
make it more complicated**

never push or pop anywhere else please

The “s” Register Terms of Service

- If you want to use an s register...
- You must **save and restore it**, just like **ra**.

my_func:

push ra

push s0

code that uses s0! it's fine! we saved it!

pop s0

pop ra

jr ra

} moving the papers off the desk

} putting the papers back

the pops happen
in **reverse order!**

Things To Consider:

- You must **always** pop the same number of registers that you push.
- To make this simpler for yourself...
make a label before the pops.
 - Then you can leave the function by jumping/branching there.
- Again: *only push at the top of the function and only pop at the bottom.*
 - Never anywhere else!

```
my_func:
    push ra
    push s0
    ...
    bge ...
    b exit_func
    ...
exit_func:
    pop s0
    pop ra
    jr ra
```