

# CS/COE 0447

Functions:  
Calling Conventions,  
and The Stack

wilkie (with content borrowed from:  
Jarrett Billingsley  
Dr. Bruce Childers)

# Lightning Recap

- We can turn if-else pseudo-code into:

```
if(s0 == 30) {
```

```
    block A
```


```
}
```

```
else {
```

```
    block B
```

```
}
```

```
block C
```



```
beq s0, 30, blockA
```

```
b blockB
```

```
blockA:
```

```
    block A
```

```
b blockC
```

```
blockB:
```

```
    block B
```

```
blockC:
```

```
    block C
```

# Our Alternative Approach

- If you see a conditional branch followed by an unconditional one...  
...you can **merge them** like so:

**invert** the condition and  
**get rid of** the first label.

```
beq s0, 30, if  
b else
```

```
bne s0, 30, else
```

if:

```
    block A
```

```
b end
```

else:

```
    block B
```

end:

```
    block A
```

```
b end
```

else:

```
    block B
```

end:


# Calling Conventions

It's kind of like computer etiquette.

# What is a Calling Convention (CS 449)

- It ensures our programs don't trip over their own feet
- It's **how machine-language functions call one another**.
  - How **arguments** are passed.
  - How **return values** are returned.
  - How control **flows into/out of the function**
  - What **contracts** exist between the caller and the callee.
  - **For instance: What registers it will use.**
- Common terms we will be using: Caller v. Callee:


```
void major() {  
    minor();  
}  
      caller
```



```
void minor() {  
    ...  
}  
      callee
```

# The Program Counter

- a program's instructions are in memory, so they have **addresses**
- the PC (program counter) holds the address of **the next instruction to run**



btw what pattern do you notice about these addresses?

# Branches and the Program Counter


- The “branch” instructions we have already seen actually simply interact with the program counter.

Instruction	Meaning
<b>beq</b> a, b, label	if(a == b) { \$pc = label }
<b>bne</b> a, b, label	if(a != b) { \$pc = label }
<b>blt</b> a, b, label	if(a < b) { \$pc = label }
<b>ble</b> a, b, label	if(a <= b) { \$pc = label }
<b>bgt</b> a, b, label	if(a > b) { \$pc = label }
<b>bge</b> a, b, label	if(a >= b) { \$pc = label }

- Each sets \$pc to the address of whatever label is given.

# Control Flow (Reprise)

- When the caller calls a function, where do we go?
- When the callee's code is finished, where do we go?





# MIPS ISA: The “jal” Instruction

- We **call** functions with **jal**: jump and link

What address should  
go into PC next?

When **func** returns,  
where will we go?

PC 0x8004

PC 0x8C30

ra 0x8008

0x8000 **li** a0, 10

0x8004 **jal** func

0x8008 **li** v0, 10

...

This is what **jal** does:

It **jumps** to a new location, and  
makes a **link** back to the old one in  
the **ra** (return address) register

**and this is ALL it does.**

**func:** 0x8C30 **li** v0, 4

...

# MIPS ISA: The “jr” Instruction

- We return from functions with **jr**: jump to address in register

Now we're at the end of **func**. **ra** still has the proper return address (*thanks to jal*)

**jr ra** copies **ra** into **pc**.

**PC** 0x8C38

**ra** 0x8008



**PC** 0x8008

0x8000 **li** a0, 10

0x8004 **jal** func

0x8008 **li** v0, 10

... ..

**func:** 0x8C30 **li** v0, 4

0x8C34 **syscall**

0x8C38 **jr** ra

and this is ALL it does.

# Don't Step on Other's Toes

- let's track **PC** and **ra** as we run this code.

	PC	ra		
	0x8000	0x0000	0x8000	<b>jal</b> fork
			0x8004	<b>li</b> v0, 10
			...	...
After <b>jal</b> fork:	0x8020	0x8004		
After <b>jal</b> spoon:	0x8040	0x8024	• fork:	0x8020 <b>jal</b> spoon
				0x8024 <b>jr</b> ra
After <b>jr</b> ra:	0x8024	0x8024		...
After <b>jr</b> ra:	0x8024	0x8024		...
After <b>jr</b> ra:	0x8024	0x8024	spoon:	0x8040 <b>jr</b> ra
After <b>jr</b> ra:	0x8024	0x8024		
After <b>jr</b> ra:	0x8024	0x8024		

UHHHHHHHHHH

# What's the Deal


- there's only **one return address register**
- if we call more than **one level deep**, things go horribly wrong
- Could we put it in another register?
  - Then what about three levels deep? four?
    - We just don't have **enough registers...**
  - **Pro-tip:** *Resist any urge to preserve registers in other registers.*
    - Just use memory.
- So **where** do we put things when we don't have room in registers?
  - Memory. But more specifically...

# The Stack

It's just memory.


# A Busy Desk

- There's a tiny desk that three people have to share
- Person 1 is working at the desk. it's covered in their stuff.
- Person 2 **interrupts them** and needs to do some important work
- **What does person 2 do with the stuff?**
  - Throw it in the trash?
- They **put it somewhere else.**
- When they are done, **they put it back.**
  - According to “etiquette” aka “calling convention” 😊




# A Very Busy Desk

- So person 2 is working...
- Then person 2 is interrupted by person 3.
- When person 3 is done, person 2 will come back.
- Where do we put person 2's stuff?
  - **on top of the stack of stuff.**
- The desk represents the **registers**.
- The people are **functions**.
- The stack of stuff is... **the stack**.



# What is the Stack? (“sp” Register)

- It's an **area of memory** provided to your program by the OS
  - When your program starts, it's already there.
- The stack is holds **information about function calls**:
  - **The return address** to the caller
  - **Copies of registers** that we want to change
  - **Local variables** that can't fit in registers
- How do we access the stack?
  - Through the **stack pointer (sp) register**
  - This register is initialized for you by the OS too.





# Accessing the Stack (animated)

- Let's say **sp** starts at the address **0xF000**
- We want to **push** something on the stack
- The first thing we'll do is **move sp to the next available slot**
- Clearly, that's the *previous* address
  - **Subtract** 4 from **sp** (why 4?)
- Then, we can store something in **the memory that sp points to.**



sp →

...	...
0xF008	0x00000000
0xF004	0x00000000
0xF000	0x00000000
0xEFFC	0xC0DEBEEF

# MIPS: Accessing the Stack (animated)

- Say `ra = 0xC0DEBEEF`
- First: move the stack pointer down (up?):

**sub** `sp, sp, 4`

- Then, store `ra` at the address that `sp` holds.

**sw** `ra, (sp)`

- Now the value in `ra` is **saved** on the stack, and **we can get it back later**.
  - And we can store as many return addresses as we want!

`sp` →

...	...
0xF008	0x00000000
0xF004	0x00000000
0xF000	0x00000000
0xEFFC	<b>0xC0DEBEEF</b>

## MIPS: Going the Other Direction (animated)

- Now we wanna **pop** the value off the stack and put it back in **ra**
- We do the same things, but in reverse

**lw** **ra**, (**sp**)

- Then, we move the stack pointer...  
up? down? whatever

**add** **sp**, **sp**, **4**

- Now we got back the old value of **ra**!
- And **sp** is back where it was before!

**ra** **0xC0DEBEEF**

...	...
0xF008	0x00000000
0xF004	0x00000000
0xF000	0x00000000
<b>sp</b> → 0xEEFC	0xC0DEBEEF

## MIPS: Simplifying: “push” and “pop” pseudo-ops

- The push and pop operations always look and work the same
- Since you'll be using them in most functions, we shortened em!
- If you write **push ra** or **pop ra**, it'll do these things for you!
  - Thank goodness.

**push** ra      =      **subi** sp, sp, 4  
                             **sw**    ra, (sp)

**pop** ra      =      **lw**    ra, (sp)  
                             **addi** sp, sp, 4

These are **pseudo-ops**: fake instructions to shorten common tasks

These can be used with ANY register, not just **ra**!

# Some Steel-Toed Boots

sp	➡	0x8004	PC	ra
			0x8000	0x0000
After <b>jal</b> fork:			0x8020	0x8004
Then we <b>push ra on the stack!</b>				
After <b>jal</b> spoon:			0x8040	0x802C
After spoon <b>jr</b> ra:			0x802C	0x802C
Then we <b>pop ra off the stack!</b>				
Before fork <b>jr</b> ra:			0x8034	0x8004
After fork <b>jr</b> ra:			0x8004	0x8004

0x8000	<b>jal</b>	fork
fork:		
0x8020	<b>push</b>	ra
0x8028	<b>jal</b>	spoon
0x802C	<b>pop</b>	ra
0x8034	<b>jr</b>	ra
spoon:		
0x8040	<b>jr</b>	ra

# Writing a Simple Function

- Writing a function follows a simple structure:

1. Give it a name (label). **spoon:**

2. Save **ra** to the stack. **push ra**

3. Do whatever. ***your code goes here***

4. Load ra from the stack. **pop ra**

5. Return! **jr ra**

- Push everything you may need. Pop it back at the end.

# Extending: Preserving other Registers

- Writing a function follows a simple structure:

1. Give it a name (label). **spoon:**

2a. Save **ra** to the stack.

**push** ra

2b. Save **s0** to the stack.

**push** s0

3. Do whatever.

*your code using s0 goes here*

4a. Load s0 from the stack.

**pop** s0

4b. Load ra from the stack.

**pop** ra

5. Return!

**jr** ra

- Push everything you may need. Pop it back **in reverse order** at the end. **(We will look at why we do this and when soon!)**

# When to “push” and “pop”

- Treat pushes and pops like the { **braces** } around a function.

**spoon:**

```
push ra # {  
# 800 instructions  
# so much stuff omg  
pop  ra # }  
jr   ra
```

pushes come at the  
beginnings of functions

pops come at the end

**That is it, seriously, don't  
make it more complicated**

***never push or pop anywhere else please***



# Calling Conventions: Arguments and Return Values

Functions that can now do things!

# Passing Arguments

- If we have a function in a higher level language...

```
v0
int gcd(a0 int a, a1 int b) {
    while(a != b) {
        if(a > b)
            a -= b;
        else
            b -= a;
    }
    return a;
}
```

We use particular registers to pass arguments and return values.

We already know how to return. How?

For this, just *put the value you want to return in **v0** before jr ra.*

# The “a” and “v” Registers

- **a0-a3** are the **argument registers**
- **v0-v1** are the **return value registers**
  - This is just a **convention**, there's nothing special about them
- To call a function...
  - You **put its arguments in the a registers** before doing a jal
- Once control is inside the callee...
  - The arguments are just "there" in the a registers.
  - Cause they are.
  - They didn't go anywhere...

# Let's Write a Function

- let's write a function like this:

```
int add_nums(int x, int y) {return x + y;}
```

- inside of our **add\_nums** asm function...
  - which register represents **x**?
  - which register represents **y**?
  - which register will hold the sum that we return?

**add\_nums:**

```
add v0, a0, a1
```

```
jr    ra
```

# Let's Call That Function

- Now let's make **main** do this:

```
print(add_nums(3, 8))
```

- How do we set 3 and 8 as the arguments?
- How do we call **add\_nums**?
- Afterwards, which register holds the sum?
- So how can we print that value?
- *Why* do syscalls put the number of the syscall in **v0**?
  - Well what do you get when you cross an elephant and a rhino?
    - Hell if I know (it's a bad, confusing decision. Sorry.)

```
li    a0, 3  
li    a1, 8  
jal   add_nums  
move  a0, v0  
li    v0, 1  
syscall
```

# Calling Conventions: Saved and Unsaved Registers

It's about how to trust functions to not step on your toes.

# An Innocent Looking Function

- Let's make a variable and a function to change it

**.data**

**counter: .word 0**

**.text**

**increment:**

**la** t0, counter

**lw** t1, (t0)

**add** t1, t1, 1

**sw** t1, (t0)

**jr** ra

then we can call it

**main:**

**jal** increment

**jal** increment

**jal** increment

# Everything's Fine Right?

- Let's write a loop that calls it ten times in a row
- So we need a loop counter ('i' in a for loop)

```
li t0, 0 # our counter
loop_begin:
    jal increment
    add t0, t0, 1
    blt t0, 10, loop_begin
loop_end:
(another way to write a for loop)
```

If we run this, it only increments the variable **once**. ☹

Why? let's put a **breakpoint** on blt and see what it sees.



## Another Piece of the Calling Convention Puzzle

- When you call a function, **it's allowed to change some registers**
- But other registers **must be left exactly as they were**

functions are required to put these registers **back the way they were before they were called.**

Saved	Unsaved
s0-s7	v0-v1
sp	a0-a3
ra*	t0-t9

anyone can change these. after you call a function, **they might have totally different values from before you called it.**

\*ra is a little weird cause it's kinda "out of sync" with the other saved regs but you DO save and restore it like the others

# When You Call a Function...

- After a **jal**, you have no idea what's in these registers.

...

**jal** increment

---

...

Unsaved

v0-v1

a0-a3

t0-t9

could be nonsense!  
garbage! bogus!

# Why It Broke

- If we look at this code again...

```
li t0, 0
loop_begin:
jal increment
add t0, t0, 1
blt t0, 10, loop_begin
loop_end:
```

**t0** is our loop counter and everything's fiiiine.

*uh oh.*

**WHAT IS IN t0 NOW??**

Instead, this is a great place to use an **s register**.

## Fixing Our Function (Abiding by Convention)

- If we use an s register...

```
li s0, 0  
loop_begin:
```

```
jal increment
```

```
add s0, s0, 1  
blt s0, 10, loop_begin
```

```
loop_end:
```

**s0** is our loop counter and everything's fiiiine.

*uh oh.*

oh whew, we used an s register, it's fine.

but s registers aren't magic. **they don't do this automatically.**

# The “s” Register Terms of Service

- If you want to use an s register...
- You must **save and restore it**, just like **ra**.

**my\_func:**

**push** ra

**push** s0

code that uses s0! it's fine! we saved it!

**pop** s0

**pop** ra

**jr** ra

} moving the papers off the desk

} putting the papers back

the pops happen  
in **reverse order!**

# Things To Consider:

- You must **always** pop the same number of registers that you push.
- To make this simpler for yourself...  
**make a label before the pops.**
  - Then you can leave the function by jumping/branching there.
- Again: *only push at the top of the function and only pop at the bottom.*
  - Never anywhere else!

```
my_func:
    push ra
    push s0
    ...
    bge ...
    b exit_func
    ...
exit_func:
    pop s0
    pop ra
    jr ra
```