

CS/COE 0447

Binary Arithmetic

wilkie (with content borrowed from:
Jarrett Billingsley
Dr. Bruce Childers)

Binary Addition

Computers and $2 + 2$ (well, $10 + 10$, eh?)

Adding in Binary

- It works the same way as you learned in school
 - Except instead of carrying at 10_{10} , you carry at... 10_2 !
 - $1 + 1 = \mathbf{10}_2$ (2_{10})
 - $1 + 1 + 1 = \mathbf{11}_2$ (3_{10})
- Let's try it. (what are these in decimal?)


$$\begin{array}{r} 1011 \ 0010 \\ +0010 \ 1111 \\ \hline \end{array}$$

Formalizing “Addition”


- For each pair of bits **starting at the LSB**,
 - Add the two bits and the carry
 - The **low bit** of the sum goes into the sum row
 - The **high bit** of the sum is the carry for the **next higher bit**
- This is the **grade school algorithm**
 - Cause it's how you learned to add in grade school
- **What if there's a carry out of the biggest column??**
 - That's ovvvvvverrrrrfloooooow (remember that?)

Ripple Adder (The Ole Classic)

- When you add one place, you might get a **carry out**.
- That becomes the **carry in** for the next higher place.



Looking at just a bit of this...



- The resulting bit for each digit is dependent on the two input bits and the carry.
- There are two outputs, however, a sum and a carry out.

1-bit Adder

- Let's try to come up with a **truth table** for adding two bits.
- Each column will hold **1 bit**.
- We will ignore the carry input, for now.

A	B	C _o	S
0	0	0	0
0	1	0	1
1	0	0	1
1	1	1	0

Half-Truth Tables

- What we just made was a **half-adder**.
- It has a carry *output* but not a carry *input*
 - (which might be useful for the lowest bit)
- To make a **full adder**, we need **3 input bits**.

C _i	A	B	C _o	S
0	0	0	0	0
0	0	1	0	1
0	1	0	0	1
0	1	1	1	0
1	0	0	0	1
1	0	1	1	0
1	1	0	1	0
1	1	1	1	1

The logic of it all


- It looks a little messy, but it kinda makes sense if you think of it like this:
 - It **counts how many input bits are "1"**
 - C_o and S are a **2-bit number!**
- If we look at the outputs in isolation:
 - S is 1 if we have an **odd number of "1s"**
 - C_o is 1 if we have **2 or 3 "1s"**
- It's a little weird, but we can build this out of AND, OR, and XOR gates

C_i	A	B	C_o	S
0	0	0	0	0
0	0	1	0	1
0	1	0	0	1
0	1	1	1	0
1	0	0	0	1
1	0	1	1	0
1	1	0	1	0
1	1	1	1	1

Boolean Expression

$$S = \overline{A}\overline{B}C_i + \overline{A}B\overline{C}_i + A\overline{B}\overline{C}_i + ABC_i$$

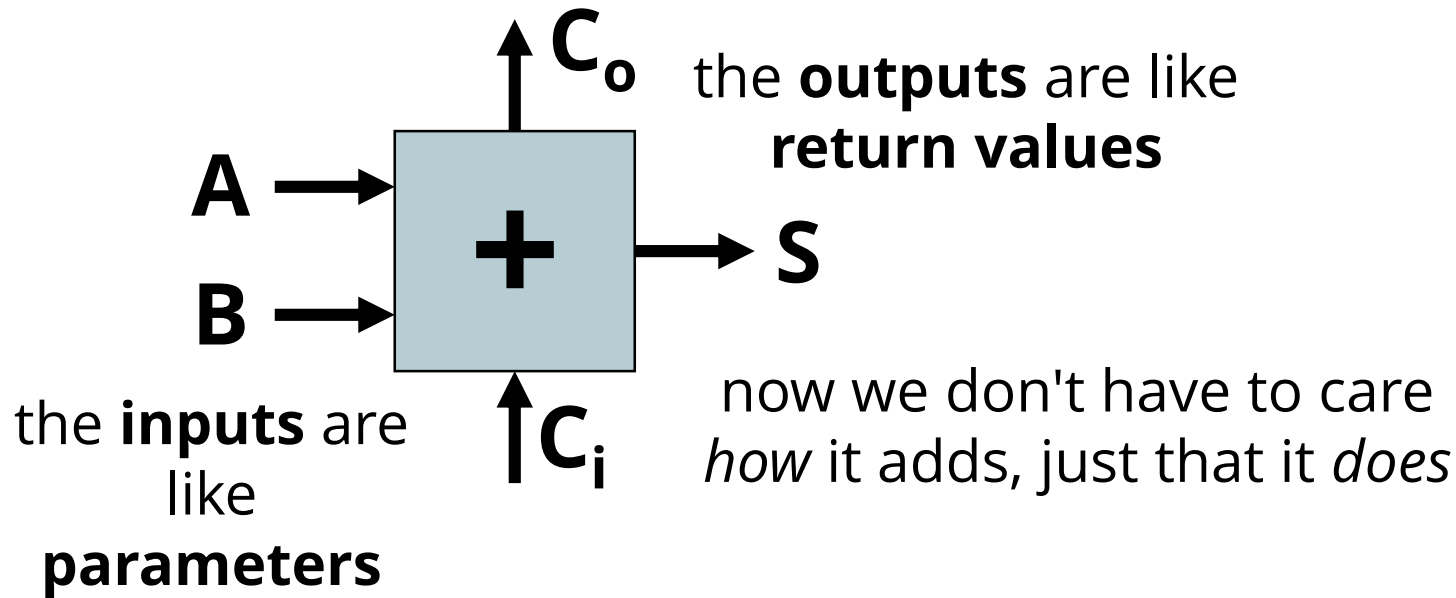
$$C_o = A\overline{B}C_i + \overline{A}B\overline{C}_i + \overline{A}BC_i + ABC_i$$



C_i	A	B	C_o	S
0	0	0	0	0
0	0	1	0	1
0	1	0	0	1
0	1	1	1	0
1	0	0	0	1
1	0	1	1	0
1	1	0	1	0
1	1	1	1	1

Sweeping it under the rug...


- In programming, we use **functions** to be able to reuse code.
- In hardware, we can group these 5 gates into a **component**.
- Here's the symbol for a **one-bit full adder**.



Adding Longer Numbers

1-bit adders are cool and all, but, like, not very useful.



Adding Longer Numbers




- The resulting bit for each digit is dependent on the two input bits and the carry.
- There are two outputs, however, a sum and a carry out.

Ripple Carry

- If we want to add two **three-bit numbers**, we'll need three **one-bit** adders.
- We chain the carries from each place to the **next higher place**, like we do on paper.
- We have to split the numbers up like so:




Negative Numbers



- That first number... is negative... hmm
- It works JUST FINE. It's really neat actually...

Recall 2's Complement




2's Complement

- Let's add them...

$$\begin{array}{rcl}
 1011 & 0010 & : -78 \\
 +0010 & 1111 & : +47 \\
 \hline
 1110 & 0001 & : -31
 \end{array}$$

Two's Complement Addition

- The great thing is: you can add numbers of either sign **without having to do anything special!**




the actual patterns of bits are the same.
so how does the computer "know" whether it's
doing signed or unsigned addition?

IT DOESN'T

What.. Even.. Is.. Subtraction?

- We *could* come up with a separate subtraction circuit, but...
- Algebra tells us that $\mathbf{x - y = x + (-y)}$
- Negation means **flip the bits** and **add 1**
- Flipping the bits uses NOT gates!
- How do we add 1 without any extra circuitry?
 - We use a *full adder* for the *LSB*, and when we're subtracting, set the "**carry in**" to **1**.




What makes a good word size?

- Can you think of an example of...
 - 100 of something?
 - a million of something? a billion?
 - a *quintillion* (10^{18})? **more?**
- $2^8 = 256$, $2^{16} \cong 65,000$, $2^{32} \cong 4$ billion, $2^{64} \cong 18$ quintillion
- For a given word size, **all the circuitry** has to be built to support it
 - 64 1-bit adders
 - 128 wires going in
 - 64 wires coming out

Gate Delay

- Electrical signals **can't move** infinitely fast
- Transistors **can't turn on and off** infinitely fast
- Since each digit must wait for the next smaller digit to compute its carry...
 - ripple carry is **linear in the number of digits**
- This is a diagram of how the outputs of a 16-bit ripple carry adder change over time
 - it's measured in *picoseconds*! so ~100ps total
- But if we went to 32 bits, it'd take 200ps
 - and 64 bits, 400ps...
- There *are* more efficient ways of adding
 - details, schmetails



Overflow

Revisiting What Happens When We Run Out of Dang Space

How many bits?


- If you add two 2-digit decimal numbers, what's the **largest number** you can get?
- What about two 4-digit decimal numbers?
- What about two 4-*bit* numbers?
- What's the pattern of the **number of digits**?
 - If you **add** two n -digit numbers *in any base*...
 - The result will have **at most $n + 1$ digits**
- That means if we add two **32-bit** numbers...
 - ...we might get a **33-bit** result!
 - (It's the 32 S output bits, and the last carry-out bit)
 - if we have more bits than we can store in our number, that's **overflow**.

$$\begin{array}{r} 99 \quad 9999 \\ +99 \quad +9999 \\ \hline 198 \quad 19998 \end{array}$$

$$\begin{array}{r} 1111 \\ +1111 \\ \hline 11110 \end{array}$$


Detecting Overflow

- For *unsigned* addition, it's easy
- For an n-bit adder:
 - just look at the C_o of the **MSB**
 - if it's 1, **it's an overflow.**
- What about subtraction?
 - shhh



Hmm, negative numbers

- The signed number line looks like this:
- Overflow occurs if we go off **either end of the number line**
- In **unsigned** 3-bit arithmetic...
 - $111_2 + 001_2 = 1000_2$
 - that is, $7 + 1 = 8$, which is **too big**.
- But in **signed** 3-bit arithmetic...
 - $111_2 + 001_2 = 0$
 - because $-1 + 1 = 0$!
- Same bit patterns, but **different results**.
 - Detecting signed overflow is a bit more subtle...



Detecting Signed Overflow

- If you add two numbers of **different signs** (negative and positive)...
 - Is it possible to **go off the ends** of the number line?
 - **no!**
 - That always gets you closer to 0
- If you add two numbers of **the same sign**, how do you know if you **went off the end of the number line**?
 - If you add two positive numbers and get a **negative number**.
 - If you add two negative numbers and get a **positive number**.
- How do you check the signs of the three numbers?
 - Do you think you could come up with a truth table? ;)

Handling overflow

- We could **ignore it**
 - In MIPS: **addu, subu**
 - This is usually a **bad idea**
 - Your program is broken
 - It's also the default in most languages, thanks C
- We could fall on the **floor - i.e. crash**
 - In MIPS **add, sub**
 - Can be handled and recovered from
 - But more complex
- We could **store that 33rd bit somewhere else**



Maybe the bit bucket is a real place...

- Many *other* architectures do store this final 33rd (etc) bit.
 - MIPS does not.
- They have a "carry bit" register
 - This can be checked by the program after an add/sub
- This is very useful for **arbitrary precision arithmetic**
 - If you want to add 2048-bit numbers, chain many 32-bit additions!